



RTX 5060 Laptop GPU · FP32 · p50 of 10-20 retained runs


(a) Attention prefill



(b) End-to-end decode



(c) One-token decode



P=128 for all cases. B=batch, D=new tokens. Random weights; systems, not quality results.